

TAUSEEF IQBAL

Machine Learning Engineer | LLM & RAG Systems

Islamabad, Pakistan | +92 346 1985132 | tauseefiqbal933@gmail.com | linkedin.com/in/tauseef-iqbal933 | github.com/tauseef933

PROFESSIONAL SUMMARY

Machine Learning Engineer with hands-on experience building production LLM and RAG systems, hybrid vector retrieval, and NLP pipelines. Raised product-mapping accuracy from ~60% to over 80% for a US-based e-commerce client at DCUBE Technologies. Proficient in Python, LangChain/LangGraph, and shipping end-to-end AI systems on free-tier and self-hosted infrastructure.

TECHNICAL SKILLS

Languages: Python, SQL, JavaScript, Dart

LLM & NLP: LangChain, LangGraph, Hugging Face Transformers, OpenAI API, Groq, RAG Systems, Hybrid Retrieval (RRF), HyDE, Prompt Engineering, AST-Aware Chunking

ML/DL & Data: TensorFlow, PyTorch, Keras, Scikit-learn, ChromaDB, Pandas, NumPy

Automation & Deployment: n8n, FastAPI, Docker, REST APIs, Git, Oracle Cloud (Always Free), Vercel, Railway

PROFESSIONAL EXPERIENCE

Junior Machine Learning Engineer — DCUBE Technologies (Pvt) Ltd, NUST Islamabad *Aug 2025 – Aug 2026*

- Raised taxonomy/field-mapping accuracy from ~60% to over 80% for a US-based e-commerce client's product mapping tool by analyzing key columns against a reference taxonomy system.
- Built an LLM-powered Promotion Optimization System (LangChain + ChromaDB) for vector search and retail insight generation, and an automated WhatsApp AI Agent for customer interactions.
- Built the Belami Asset Automation Tool (Streamlit) to automate enterprise asset template processing and validation for vendor Excel files; managed Akeneo product data and image-processing pipelines.

Machine Learning Engineer Intern — DCUBE Technologies (Pvt) Ltd, NUST Islamabad *Apr 2025 – July 2025*

- Engineered a retail promotion optimization system with Gradio and Flutter UI; fine-tuned Hugging Face transformer models for personalized recommendations and uplift prediction.

KEY PROJECTS

CodeSage — Natural-Language Code Search (RAG) — *Python, LangChain, Hybrid Retrieval, HyDE, React*

- Built a RAG system that answers natural-language questions about a codebase and cites exact file/line evidence, using AST-aware chunking, hybrid retrieval with RRF, and HyDE to boost retrieval precision over naive vector search.
- Designed a Groq/Gemini LLM router balancing cost, latency, and quality; shipped a React frontend, deployed on Oracle Cloud Always Free.

LexMind AI — Multi-Agent Legal RAG Platform — *LangGraph, Groq (Llama 3.1 70B), ChromaDB, FastAPI, React*

- Built a multi-agent legal intelligence platform using LangGraph to orchestrate specialized agents over a ChromaDB vector store for legal document retrieval, risk detection, and timeline analysis; frontend live on Vercel.

NEXUS — Multi-Agent Voice Assistant — *Python, FastAPI, Groq Whisper/Llama 3.3-70B, WebSocket*

- Built a CPU-only voice assistant that transcribes speech and routes each turn to one of four specialist agents, runs tools, and replies via neural TTS, with an Ollama fallback LLM for rate-limit resilience.
- Supports bilingual conversations in English and Roman Urdu, including transliteration handling to fix Roman Urdu pronunciation in the TTS output.

EDUCATION & CERTIFICATIONS

B.S. Computer Science — COMSATS University Islamabad, Abbottabad Campus (2021–2025), Specialization: AI, ML, Mobile Development

Certifications: Machine Learning, Advanced Learning Algorithms, NLP, AI For Everyone (Stanford University); ChatGPT Prompt Engineering (DeepLearning.AI); UX Design Foundations (Google)